

Keeping track of things

Price point \$3.5k for 16 DBs.

Meeting 1 Feedback: 📅 Wednesday, February 11th

Okay talked with Spencer —

- He said client will send SOW later after finishing this initial 14
- To just rework the SQL descriptions for these existing ones until then
- Client is super super super important because this is the only thing we're doing for them at the moment, so Spencer's willing to pay your \$3k-4k ish average db rate if this scales up, even if this is at a loss to us so please don't worry about that part !
- Client doesn't want complexity, quality to diminish as it scales up

For the queries

- He just emphasized human-made rewrite/edits and that the descriptions of the queries should be more detailed about the intent of the query (not a summary)
- Also looks like some of the jsons don't match the .md in terms of description and detail so they might have been talking about that too
- — (db-2, db-3, db-6, db-7, db-10 json files were missing intent/purpose fields)
- **Have longer descriptions per query (more deeply contextualized)**

Pasting from Spencer

Feedback from client

Databases look good, only comment: the descriptions are AI-generated

- **To-do: Need to go in and make the descriptions all 100% human made**

- **To-do: Query descriptions: should be updated to talk about query intent / what they are aiming to do rather than a summary of the query**
- Next steps
- Client to send over a file format to ingest data
- Client concerned about us increasing complexity of the databases / queries as we scale up
- **Have longer descriptions per query (more deeply contextualized)**

Meeting 2: 📅 Wednesday, February 18th

- **Client Meeting Prep Sheet *For Spencer* – compiled by Alan**
Open with alignment
- "Before we dive in, can you walk us through what your ideal final deliverable looks like end to end?"
- "Is there a written rubric or spec we can reference for quality reviews going forward?"
- **Requirements**
- What defines a 5/5 query description versus a 1/5? Can you show us an example of each?
- Is the primary concern SQL correctness, natural language quality, or both?
- When you flag something as AI-generated – what specifically does that mean to you? Generic language, missing context, something else?
- Are you comparing this dataset against an existing benchmark like BIRD or Spider? What methodologies will you use to test it?
- **Process**
- Who is the final sign-off on delivery acceptance – is that you or someone else on your team?
- What does your internal review process look like – who reviews, what tool, in what order?
- How much lead time do you need to review a delivery before a Wednesday meeting?
- **Success criteria**
- What is your target execution accuracy or F1 score?

- How many queries per batch and what's the expected turnaround?
- What would make you confident enough to scale this engagement?
- **Close with a boundary (friendly)**
- "We have a senior expert on this who is operating at a level that's hard to find — to protect turnaround time and quality, it would really help us to lock in a written spec after this call so she can execute cleanly without mid-cycle pivots. Can we do that today?"
- **Upsell — Future Opportunity**
Plant the seed, don't pitch hard. The goal is to get them thinking bigger.
- "As we're locking in quality standards here — are you thinking about scaling this dataset beyond the current batch?"
- "What does your roadmap look like for the data agent — are you at the build phase or still validating?"
- "One thing worth knowing is that the expert we have on this built the entire pipeline from scratch including the infrastructure, SQL logic, and benchmark alignment. Most vendors can't offer that in one person. If you're thinking about scaling, that's a significant advantage to have on your side."
- "If quality and consistency are the priority as you scale — we can build a rubric together now that becomes the standard for every future batch. That protects you and makes reviews faster on your end."

<https://bird-bench.github.io/>

MEDDPICC - Clearly client is not happy with you.

LLMs cannot determine what good quality data that can be used in benchmarks or training or substitute actual ML/DS expertise with shipping an actual product like an multilingual weather translation model and cannot provide the actual metrics of ML/DS workloads that indicate the data is good for training, benchmark, etc or developing

	any product like their terminus agent.
--	--